

```
In [ ]: import pandas as pd
import numpy as np
import os
import shutil
import random
from tqdm import tqdm
from pathlib import Path
import cv2 as cv
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras import layers, optimizers, losses, metrics, callbacks
from tensorflow.keras import Sequential, Model, Input
from sklearn.model_selection import train_test_split
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from sklearn.model_selection import train_test_split
import warnings
warnings.filterwarnings("ignore")
random.seed(45)

print(tf.__version__)
```

2.19.0

```
In [ ]: import glob
from PIL import Image, ImageFile
from joblib import Parallel, delayed
ImageFile.LOAD_TRUNCATED_IMAGES = True
```

## LOAD IMAGE DATA AND UNDERSTANDING SOME PROPERTIES OF IMAGE

```
In [3]: def image_properties(path):
    for img in random.sample(os.listdir(path),1):
        print('Image name =',img)
        image = cv.imread(os.path.join(path, img),cv.IMREAD_COLOR)
        break


    return image
```

```
In [4]: HOME_PATH = os.getcwd() + "/"

path = HOME_PATH + 'train_images/normal'

image = image_properties(path)
plt.imshow(image)
plt.title('Normal image')
print(f"The dimensions are {image.shape[0]} pixels height and {image.shape[1]} pixels width")
print(f"The maximum pixel value is {image.max():.2f}")
print(f"The minimum pixel value is {image.min():.2f}")
print(f"The mean value of the pixels is {image.mean():.2f}")
print(f"The standard deviation is {image.std():.2f}")
```

Image name = 103292.jpg  
The dimensions are 640 pixels height and 480 pixels width  
The maximum pixel value is 255.00  
The minimum pixel value is 0.00  
The mean value of the pixels is 120.33  
The standard deviation is 69.71



### Loading a dataset

```
In [5]: batch_size = 64
img_height = 256
img_width = 256

data_dir = HOME_PATH + 'train_images'
```

### Image Data-Generator

### Data Normalization And Data Augmentation

```
In [6]: img_datagen = tf.keras.preprocessing.image.ImageDataGenerator(
    rescale=1.0/255.0,
    validation_split=0.3,
    rotation_range=5,
    shear_range=0.3,
    zoom_range=0.3,
    width_shift_range=0.05,
    height_shift_range=0.05,
    horizontal_flip=True,
```

```
) vertical_flip=True
```

### Training Dataset (70%)

```
In [7]: train_gen = img_datagen.flow_from_directory(
        data_dir,
        subset="training",
        seed=42,
        target_size=(img_height, img_width),
        batch_size=batch_size,
        class_mode="categorical",
        color_mode='rgb'
    )
```

Found 7288 images belonging to 10 classes.

### Testing Dataset (30%)

```
In [8]: valid_gen = img_datagen.flow_from_directory(
        data_dir,
        subset="validation",
        seed=42,
        target_size=(img_height, img_width),
        batch_size=batch_size,
        class_mode="categorical"
    )
```

Found 3119 images belonging to 10 classes.

### Class Lables

```
In [9]: print('Total No Of Classes in the datasetL:', len(train_gen.class_indices))
        print('Class Names:', train_gen.class_indices)
```

Total No Of Classes in the datasetL: 10

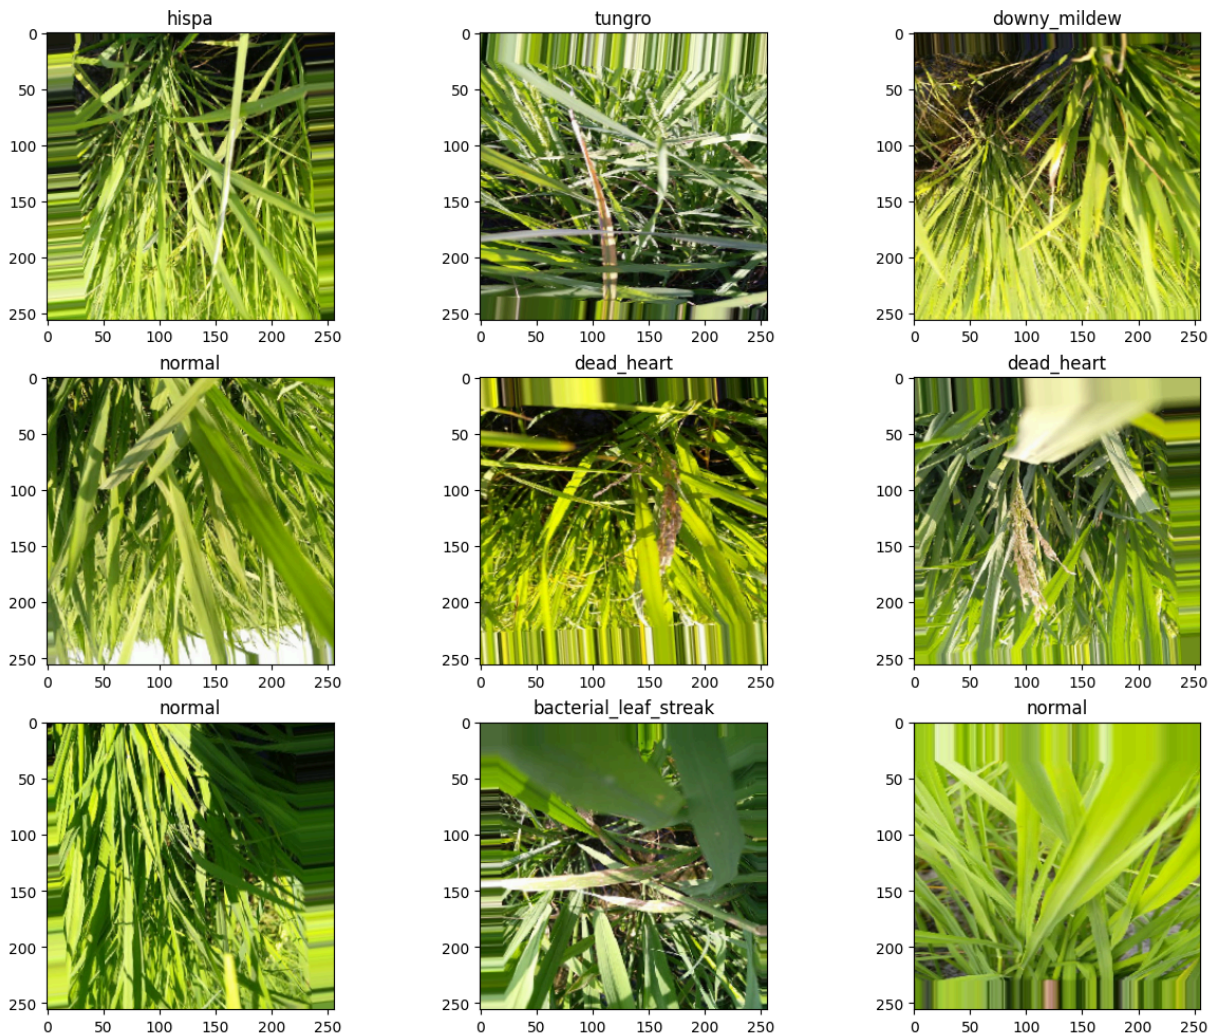
Class Names: {'bacterial\_leaf\_blight': 0, 'bacterial\_leaf\_streak': 1, 'bacterial\_panicle\_blight': 2, 'blast': 3, 'brown\_spot': 4, 'dead\_heart': 5, 'downy\_mildew': 6, 'hispal': 7, 'normal': 8, 'tungro': 9}

### Displaying a 9 Random images form dataset

```
In [10]: fig, axs = plt.subplots(nrows=3, ncols=3, figsize=(15, 12))
        plt.subplots_adjust(hspace=0.2)
        fig.suptitle("9 Random images form dataset", fontsize=18, y=0.95)

        for i in range(1,10):
            plt.subplot(3,3,i)
            img, label = train_gen.__next__()
            plt.title(list(train_gen.class_indices.keys())[np.argmax(label)])
            plt.imshow(img[0])
```

## 9 Random images form dataset



## Vision Transformer (ViT)

REF:

[https://keras.io/examples/vision/image\\_classification\\_with\\_vision\\_transformer/](https://keras.io/examples/vision/image_classification_with_vision_transformer/)

```
In [11]: # Loads an image and convert them to numpy array
def load_images(paths):
    data = []
    labels = []
    i = 0
    for label, path in tqdm(enumerate(paths)):
        for img_path in os.listdir(path):
            image = np.array(Image.open(os.path.join(path, img_path)).convert('RGB').resize(
                data.append(image)
                labels.append(label)

    return np.array(data), np.asarray(labels)
```

```
In [12]: images, labels = load_images(glob.glob(HOME_PATH + 'train_images/*'))
```

```
10it [00:52, 5.23s/it]
```

### Train Test Split (70:30)

```
In [13]: X_train, X_test, y_train, y_test = train_test_split(images, labels.reshape(-1,1), t
```

```
In [14]: num_classes = 10  
input_shape = (256, 256, 3)
```

### Configure the hyperparameters

```
In [15]: learning_rate = 0.001  
weight_decay = 0.0001  
batch_size = 32  
num_epochs = 100  
image_size = 72  
patch_size = 6  
num_patches = (image_size // patch_size) ** 2  
projection_dim = 64  
num_heads = 4  
transformer_units = [  
    projection_dim * 2,  
    projection_dim,  
] # Size of the transformer layers  
transformer_layers = 8  
mlp_head_units = [2048, 1024] # Size of the dense layers of the final classifier
```

### Data Augmentation

```
In [16]: data_augmentation = keras.Sequential(  
    [  
        layers.Normalization(),  
        layers.Resizing(image_size, image_size),  
        layers.RandomFlip("horizontal"),  
        layers.RandomRotation(factor=0.02),  
        layers.RandomZoom(  
            height_factor=0.2, width_factor=0.2  
        ),  
    ],  
    name="data_augmentation",  
)  
# Compute the mean and the variance of the training data for normalization.  
data_augmentation.layers[0].adapt(X_train)
```

## Multilayer perceptron (MLP)

```
In [17]: def mlp(x, hidden_units, dropout_rate):  
    for units in hidden_units:  
        x = layers.Dense(units, activation=tf.nn.gelu)(x)  
        x = layers.Dropout(dropout_rate)(x)  
    return x
```

## Implementing patch creation as a layer

```
In [18]: class Patches(layers.Layer):
    def __init__(self, patch_size):
        super(Patches, self).__init__()
        self.patch_size = patch_size

    def call(self, images):
        batch_size = tf.shape(images)[0]
        patches = tf.image.extract_patches(
            images=images,
            sizes=[1, self.patch_size, self.patch_size, 1],
            strides=[1, self.patch_size, self.patch_size, 1],
            rates=[1, 1, 1, 1],
            padding="VALID",
        )
        patch_dims = patches.shape[-1]
        patches = tf.reshape(patches, [batch_size, -1, patch_dims])
        return patches
```


```
In [19]: plt.figure(figsize=(4, 4))
image = X_train[np.random.choice(range(X_train.shape[0]))]
plt.imshow(image.astype("uint8"))
plt.axis("off")

resized_image = tf.image.resize(
    tf.convert_to_tensor([image]), size=(image_size, image_size)
)
patches = Patches(patch_size)(resized_image)
print(f"Image size: {image_size} X {image_size}")
print(f"Patch size: {patch_size} X {patch_size}")
print(f"Patches per image: {patches.shape[1]}")
print(f"Elements per patch: {patches.shape[-1]}")

n = int(np.sqrt(patches.shape[1]))
plt.figure(figsize=(4, 4))
for i, patch in enumerate(patches[0]):
    ax = plt.subplot(n, n, i + 1)
    patch_img = tf.reshape(patch, (patch_size, patch_size, 3))
    plt.imshow(patch_img.numpy().astype("uint8"))
    plt.axis("off")
```

Image size: 72 X 72  
Patch size: 6 X 6  
Patches per image: 144  
Elements per patch: 108





### Implementing the patch encoding layer

```
In [20]: class PatchEncoder(layers.Layer):
def __init__(self, num_patches, projection_dim):
    super(PatchEncoder, self).__init__()
    self.num_patches = num_patches
    self.projection = layers.Dense(units=projection_dim)
    self.position_embedding = layers.Embedding(
        input_dim=num_patches, output_dim=projection_dim
    )

def call(self, patch):
    positions = tf.range(start=0, limit=self.num_patches, delta=1)
    encoded = self.projection(patch) + self.position_embedding(positions)
    return encoded
```

## ViT Model

```
In [21]: def create_vit_classifier():

    inputs = layers.Input(shape=input_shape)
    # Augment data.
    augmented = data_augmentation(inputs)
    # Create patches.
    patches = Patches(patch_size)(augmented)
    # Encode patches.
    encoded_patches = PatchEncoder(num_patches, projection_dim)(patches)

    # Create multiple layers of the Transformer block.
    for _ in range(transformer_layers):
        # Layer normalization 1.
        x1 = layers.LayerNormalization(epsilon=1e-6)(encoded_patches)
        # Create a multi-head attention layer.
        attention_output = layers.MultiHeadAttention(
            num_heads=num_heads, key_dim=projection_dim, dropout=0.1
        )(x1, x1)
        # Skip connection 1.
        x2 = layers.Add()([attention_output, encoded_patches])
        # Layer normalization 2.
        x3 = layers.LayerNormalization(epsilon=1e-6)(x2)
        # MLP.
        x3 = mlp(x3, hidden_units=transformer_units, dropout_rate=0.1)
        # Skip connection 2.
        encoded_patches = layers.Add()([x3, x2])

    # Create a [batch_size, projection_dim] tensor.
    representation = layers.LayerNormalization(epsilon=1e-6)(encoded_patches)
    representation = layers.Flatten()(representation)
    representation = layers.Dropout(0.5)(representation)
    # Add MLP.
    features = mlp(representation, hidden_units=mlp_head_units, dropout_rate=0.5)
    # Classify outputs.
    logits = layers.Dense(num_classes)(features)
    # Create the Keras model.
    model = keras.Model(inputs=inputs, outputs=logits)
    return model
```

## Callback

```
In [27]: filepath = HOME_PATH + 'paddy_models/model_vgg_new.keras'

checkpoint = tf.keras.callbacks.ModelCheckpoint(
    filepath = filepath,
    monitor="val_accuracy",
    verbose=1,
    save_best_only=True,
    mode = 'auto'
)
```



```
In [28]: class TerminateNaN(tf.keras.callbacks.Callback):
def on_epoch_end(self, epoch, logs={}):
    loss = logs.get('loss')
    if loss is not None:
        if np.isnan(loss) or np.isinf(loss):
            print('Invalid loss and terminated at loss {}'.format(epoch))
            self.model.stop_training = True

terminate_nan = TerminateNaN()
```

```
In [29]: # Learning rate scheduler
reduce_lr = tf.keras.callbacks.ReduceLROnPlateau(
    monitor="val_accuracy",
    factor=0.1,
    patience=10,
    verbose=1,
    mode="auto",
    min_delta=0.001,
    cooldown=3,
    min_lr=0
)
```

```
In [30]: def train(model):
optimizer = keras.optimizers.AdamW(
    learning_rate=learning_rate, weight_decay=weight_decay
)

model.compile(
    optimizer=optimizer,
    loss=keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=[
        keras.metrics.SparseCategoricalAccuracy(name="accuracy"),
    ]
)

history = model.fit(
    x=X_train,
    y=y_train,
    batch_size=batch_size,
    epochs=num_epochs,
    validation_split=0.3,
    callbacks=[checkpoint, reduce_lr, terminate_nan]
)

_, accuracy = model.evaluate(X_test, y_test)
print(f"Test accuracy: {round(accuracy * 100, 2)}%")

return history

vit_classifier = create_vit_classifier()
vit_classifier.summary()
```

Model: "functional\_2"

Layer (type)	Output Shape	Param #	Connected to
input_layer_2 (InputLayer)	(None, 256, 256, 3)	0	-
data_augmentation (Sequential)	(None, 72, 72, 3)	7	input_layer_2[0]...
patches_2 (Patches)	(None, None, 108)	0	data_augmentatio...
patch_encoder_1 (PatchEncoder)	(None, 144, 64)	16,192	patches_2[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	patch_encoder_1[...
multi_head_attenti... (MultiHeadAttentio...	(None, 144, 64)	66,368	layer_normalizat... layer_normalizat...
add_16 (Add)	(None, 144, 64)	0	multi_head_atten... patch_encoder_1[...
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_16[0][0]
dense_21 (Dense)	(None, 144, 128)	8,320	layer_normalizat...
dropout_28 (Dropout)	(None, 144, 128)	0	dense_21[0][0]
dense_22 (Dense)	(None, 144, 64)	8,256	dropout_28[0][0]
dropout_29 (Dropout)	(None, 144, 64)	0	dense_22[0][0]
add_17 (Add)	(None, 144, 64)	0	dropout_29[0][0], add_16[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_17[0][0]
multi_head_attenti... (MultiHeadAttentio...	(None, 144, 64)	66,368	layer_normalizat... layer_normalizat...
add_18 (Add)	(None, 144, 64)	0	multi_head_atten... add_17[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_18[0][0]
dense_23 (Dense)	(None, 144, 128)	8,320	layer_normalizat...
dropout_31 (Dropout)	(None, 144, 128)	0	dense_23[0][0]

dense_24 (Dense)	(None, 144, 64)	8,256	dropout_31[0][0]
dropout_32 (Dropout)	(None, 144, 64)	0	dense_24[0][0]
add_19 (Add)	(None, 144, 64)	0	dropout_32[0][0], add_18[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_19[0][0]
multi_head_attenti... (MultiHeadAttentio...	(None, 144, 64)	66,368	layer_normalizat... layer_normalizat...
add_20 (Add)	(None, 144, 64)	0	multi_head_atten... add_19[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_20[0][0]
dense_25 (Dense)	(None, 144, 128)	8,320	layer_normalizat...
dropout_34 (Dropout)	(None, 144, 128)	0	dense_25[0][0]
dense_26 (Dense)	(None, 144, 64)	8,256	dropout_34[0][0]
dropout_35 (Dropout)	(None, 144, 64)	0	dense_26[0][0]
add_21 (Add)	(None, 144, 64)	0	dropout_35[0][0], add_20[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_21[0][0]
multi_head_attenti... (MultiHeadAttentio...	(None, 144, 64)	66,368	layer_normalizat... layer_normalizat...
add_22 (Add)	(None, 144, 64)	0	multi_head_atten... add_21[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_22[0][0]
dense_27 (Dense)	(None, 144, 128)	8,320	layer_normalizat...
dropout_37 (Dropout)	(None, 144, 128)	0	dense_27[0][0]
dense_28 (Dense)	(None, 144, 64)	8,256	dropout_37[0][0]
dropout_38 (Dropout)	(None, 144, 64)	0	dense_28[0][0]
add_23 (Add)	(None, 144, 64)	0	dropout_38[0][0],

			add_22[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_23[0][0]
multi_head_attenti... (MultiHeadAttentio...	(None, 144, 64)	66,368	layer_normalizatio... layer_normalizatio...
add_24 (Add)	(None, 144, 64)	0	multi_head_atten... add_23[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_24[0][0]
dense_29 (Dense)	(None, 144, 128)	8,320	layer_normalizatio...
dropout_40 (Dropout)	(None, 144, 128)	0	dense_29[0][0]
dense_30 (Dense)	(None, 144, 64)	8,256	dropout_40[0][0]
dropout_41 (Dropout)	(None, 144, 64)	0	dense_30[0][0]
add_25 (Add)	(None, 144, 64)	0	dropout_41[0][0], add_24[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_25[0][0]
multi_head_attenti... (MultiHeadAttentio...	(None, 144, 64)	66,368	layer_normalizatio... layer_normalizatio...
add_26 (Add)	(None, 144, 64)	0	multi_head_atten... add_25[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_26[0][0]
dense_31 (Dense)	(None, 144, 128)	8,320	layer_normalizatio...
dropout_43 (Dropout)	(None, 144, 128)	0	dense_31[0][0]
dense_32 (Dense)	(None, 144, 64)	8,256	dropout_43[0][0]
dropout_44 (Dropout)	(None, 144, 64)	0	dense_32[0][0]
add_27 (Add)	(None, 144, 64)	0	dropout_44[0][0], add_26[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_27[0][0]
multi_head_attenti...	(None, 144, 64)	66,368	layer_normalizatio...

(MultiHeadAttentio...			layer_normalizat...
add_28 (Add)	(None, 144, 64)	0	multi_head_atten... add_27[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_28[0][0]
dense_33 (Dense)	(None, 144, 128)	8,320	layer_normalizat...
dropout_46 (Dropout)	(None, 144, 128)	0	dense_33[0][0]
dense_34 (Dense)	(None, 144, 64)	8,256	dropout_46[0][0]
dropout_47 (Dropout)	(None, 144, 64)	0	dense_34[0][0]
add_29 (Add)	(None, 144, 64)	0	dropout_47[0][0], add_28[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_29[0][0]
multi_head_attenti... (MultiHeadAttentio...	(None, 144, 64)	66,368	layer_normalizat... layer_normalizat...
add_30 (Add)	(None, 144, 64)	0	multi_head_atten... add_29[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_30[0][0]
dense_35 (Dense)	(None, 144, 128)	8,320	layer_normalizat...
dropout_49 (Dropout)	(None, 144, 128)	0	dense_35[0][0]
dense_36 (Dense)	(None, 144, 64)	8,256	dropout_49[0][0]
dropout_50 (Dropout)	(None, 144, 64)	0	dense_36[0][0]
add_31 (Add)	(None, 144, 64)	0	dropout_50[0][0], add_30[0][0]
layer_normalizatio... (LayerNormalizatio...	(None, 144, 64)	128	add_31[0][0]
flatten_1 (Flatten)	(None, 9216)	0	layer_normalizat...
dropout_51 (Dropout)	(None, 9216)	0	flatten_1[0][0]
dense_37 (Dense)	(None, 2048)	18,876,416	dropout_51[0][0]

dropout_52 (Dropout)	(None, 2048)	0	dense_37[0][0]
dense_38 (Dense)	(None, 1024)	2,098,176	dropout_52[0][0]
dropout_53 (Dropout)	(None, 1024)	0	dense_38[0][0]
dense_39 (Dense)	(None, 10)	10,250	dropout_53[0][0]

**Total params:** 21,666,769 (82.65 MB)

**Trainable params:** 21,666,762 (82.65 MB)

**Non-trainable params:** 7 (32.00 B)

In [31]: `history = train(vit_classifier)`



Epoch 1/100  
153/153 ————— 0s 557ms/step - accuracy: 0.1647 - loss: 3.9922  
Epoch 1: val\_accuracy improved from -inf to 0.26673, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 122s 642ms/step - accuracy: 0.1649 - loss: 3.9853 - val\_accuracy: 0.2667 - val\_loss: 2.0816 - learning\_rate: 0.0010  
Epoch 2/100  
153/153 ————— 0s 560ms/step - accuracy: 0.2501 - loss: 2.1435  
Epoch 2: val\_accuracy improved from 0.26673 to 0.33174, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 627ms/step - accuracy: 0.2502 - loss: 2.1433 - val\_accuracy: 0.3317 - val\_loss: 1.9000 - learning\_rate: 0.0010  
Epoch 3/100  
153/153 ————— 0s 608ms/step - accuracy: 0.2767 - loss: 2.0216  
Epoch 3: val\_accuracy improved from 0.33174 to 0.34130, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 103s 676ms/step - accuracy: 0.2767 - loss: 2.0217 - val\_accuracy: 0.3413 - val\_loss: 1.9018 - learning\_rate: 0.0010  
Epoch 4/100  
153/153 ————— 0s 557ms/step - accuracy: 0.3043 - loss: 1.9498  
Epoch 4: val\_accuracy improved from 0.34130 to 0.34704, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 625ms/step - accuracy: 0.3043 - loss: 1.9499 - val\_accuracy: 0.3470 - val\_loss: 1.8529 - learning\_rate: 0.0010  
Epoch 5/100  
153/153 ————— 0s 556ms/step - accuracy: 0.3334 - loss: 1.9343  
Epoch 5: val\_accuracy improved from 0.34704 to 0.39771, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 95s 624ms/step - accuracy: 0.3333 - loss: 1.9341 - val\_accuracy: 0.3977 - val\_loss: 1.7562 - learning\_rate: 0.0010  
Epoch 6/100  
153/153 ————— 0s 557ms/step - accuracy: 0.3544 - loss: 1.8274  
Epoch 6: val\_accuracy improved from 0.39771 to 0.42017, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 627ms/step - accuracy: 0.3545 - loss: 1.8273 - val\_accuracy: 0.4202 - val\_loss: 1.6521 - learning\_rate: 0.0010  
Epoch 7/100  
153/153 ————— 0s 558ms/step - accuracy: 0.3890 - loss: 1.7772  
Epoch 7: val\_accuracy improved from 0.42017 to 0.44790, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 625ms/step - accuracy: 0.3890 - loss: 1.7771 - val\_accuracy: 0.4479 - val\_loss: 1.6064 - learning\_rate: 0.0010  
Epoch 8/100  
153/153 ————— 0s 557ms/step - accuracy: 0.4131 - loss: 1.7098  
Epoch 8: val\_accuracy improved from 0.44790 to 0.50335, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 95s 624ms/step - accuracy: 0.4131 - loss: 1.7098 - val\_accuracy: 0.5033 - val\_loss: 1.4756 - learning\_rate: 0.0010  
Epoch 9/100  
153/153 ————— 0s 557ms/step - accuracy: 0.4257 - loss: 1.6820  
Epoch 9: val\_accuracy improved from 0.50335 to 0.51673, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 625ms/step - accuracy: 0.4258 - loss: 1.6816 - val\_accuracy: 0.5167 - val\_loss: 1.4301 - learning\_rate: 0.0010  
Epoch 10/100  
153/153 ————— 0s 557ms/step - accuracy: 0.4571 - loss: 1.5784

Epoch 10: val\_accuracy improved from 0.51673 to 0.53107, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 625ms/step - accuracy: 0.4572 - loss: 1.5782 - val\_accuracy: 0.5311 - val\_loss: 1.3678 - learning\_rate: 0.0010  
Epoch 11/100  
153/153 ————— 0s 557ms/step - accuracy: 0.4878 - loss: 1.5049  
Epoch 11: val\_accuracy improved from 0.53107 to 0.58031, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 625ms/step - accuracy: 0.4878 - loss: 1.5047 - val\_accuracy: 0.5803 - val\_loss: 1.2611 - learning\_rate: 0.0010  
Epoch 12/100  
153/153 ————— 0s 557ms/step - accuracy: 0.4870 - loss: 1.4878  
Epoch 12: val\_accuracy improved from 0.58031 to 0.59799, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 625ms/step - accuracy: 0.4871 - loss: 1.4876 - val\_accuracy: 0.5980 - val\_loss: 1.2277 - learning\_rate: 0.0010  
Epoch 13/100  
153/153 ————— 0s 557ms/step - accuracy: 0.5385 - loss: 1.3913  
Epoch 13: val\_accuracy improved from 0.59799 to 0.60421, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 625ms/step - accuracy: 0.5385 - loss: 1.3911 - val\_accuracy: 0.6042 - val\_loss: 1.2028 - learning\_rate: 0.0010  
Epoch 14/100  
153/153 ————— 0s 558ms/step - accuracy: 0.5382 - loss: 1.3816  
Epoch 14: val\_accuracy improved from 0.60421 to 0.61950, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 626ms/step - accuracy: 0.5382 - loss: 1.3815 - val\_accuracy: 0.6195 - val\_loss: 1.1832 - learning\_rate: 0.0010  
Epoch 15/100  
153/153 ————— 0s 558ms/step - accuracy: 0.5644 - loss: 1.2833  
Epoch 15: val\_accuracy improved from 0.61950 to 0.63289, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 626ms/step - accuracy: 0.5644 - loss: 1.2834 - val\_accuracy: 0.6329 - val\_loss: 1.1070 - learning\_rate: 0.0010  
Epoch 16/100  
153/153 ————— 0s 557ms/step - accuracy: 0.5784 - loss: 1.2389  
Epoch 16: val\_accuracy improved from 0.63289 to 0.67065, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 625ms/step - accuracy: 0.5784 - loss: 1.2388 - val\_accuracy: 0.6707 - val\_loss: 0.9989 - learning\_rate: 0.0010  
Epoch 17/100  
153/153 ————— 0s 550ms/step - accuracy: 0.6073 - loss: 1.1956  
Epoch 17: val\_accuracy improved from 0.67065 to 0.68212, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 94s 618ms/step - accuracy: 0.6074 - loss: 1.1956 - val\_accuracy: 0.6821 - val\_loss: 0.9817 - learning\_rate: 0.0010  
Epoch 18/100  
153/153 ————— 0s 550ms/step - accuracy: 0.6195 - loss: 1.1429  
Epoch 18: val\_accuracy did not improve from 0.68212  
153/153 ————— 93s 611ms/step - accuracy: 0.6195 - loss: 1.1431 - val\_accuracy: 0.6625 - val\_loss: 1.0301 - learning\_rate: 0.0010  
Epoch 19/100  
153/153 ————— 0s 552ms/step - accuracy: 0.6357 - loss: 1.1084  
Epoch 19: val\_accuracy improved from 0.68212 to 0.71558, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 95s 619ms/step - accuracy: 0.6357 - loss: 1.1082 - val\_

accuracy: 0.7156 - val\_loss: 0.8844 - learning\_rate: 0.0010  
Epoch 20/100  
153/153 ————— 0s 551ms/step - accuracy: 0.6640 - loss: 0.9959  
Epoch 20: val\_accuracy improved from 0.71558 to 0.72467, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 95s 619ms/step - accuracy: 0.6640 - loss: 0.9960 - val\_accuracy: 0.7247 - val\_loss: 0.8535 - learning\_rate: 0.0010  
Epoch 21/100  
153/153 ————— 0s 552ms/step - accuracy: 0.6618 - loss: 1.0185  
Epoch 21: val\_accuracy improved from 0.72467 to 0.76243, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 95s 619ms/step - accuracy: 0.6618 - loss: 1.0184 - val\_accuracy: 0.7624 - val\_loss: 0.7804 - learning\_rate: 0.0010  
Epoch 22/100  
153/153 ————— 0s 552ms/step - accuracy: 0.6860 - loss: 0.9509  
Epoch 22: val\_accuracy did not improve from 0.76243  
153/153 ————— 94s 612ms/step - accuracy: 0.6860 - loss: 0.9509 - val\_accuracy: 0.7615 - val\_loss: 0.7812 - learning\_rate: 0.0010  
Epoch 23/100  
153/153 ————— 0s 552ms/step - accuracy: 0.6971 - loss: 0.9114  
Epoch 23: val\_accuracy did not improve from 0.76243  
153/153 ————— 94s 612ms/step - accuracy: 0.6971 - loss: 0.9115 - val\_accuracy: 0.7395 - val\_loss: 0.8377 - learning\_rate: 0.0010  
Epoch 24/100  
153/153 ————— 0s 553ms/step - accuracy: 0.7016 - loss: 0.9146  
Epoch 24: val\_accuracy improved from 0.76243 to 0.79302, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 95s 622ms/step - accuracy: 0.7017 - loss: 0.9143 - val\_accuracy: 0.7930 - val\_loss: 0.6699 - learning\_rate: 0.0010  
Epoch 25/100  
153/153 ————— 0s 552ms/step - accuracy: 0.7408 - loss: 0.8196  
Epoch 25: val\_accuracy did not improve from 0.79302  
153/153 ————— 94s 613ms/step - accuracy: 0.7407 - loss: 0.8198 - val\_accuracy: 0.7868 - val\_loss: 0.6967 - learning\_rate: 0.0010  
Epoch 26/100  
153/153 ————— 0s 554ms/step - accuracy: 0.7341 - loss: 0.7888  
Epoch 26: val\_accuracy improved from 0.79302 to 0.80497, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 95s 622ms/step - accuracy: 0.7342 - loss: 0.7888 - val\_accuracy: 0.8050 - val\_loss: 0.6257 - learning\_rate: 0.0010  
Epoch 27/100  
153/153 ————— 0s 553ms/step - accuracy: 0.7599 - loss: 0.7497  
Epoch 27: val\_accuracy did not improve from 0.80497  
153/153 ————— 94s 615ms/step - accuracy: 0.7598 - loss: 0.7498 - val\_accuracy: 0.7983 - val\_loss: 0.6462 - learning\_rate: 0.0010  
Epoch 28/100  
153/153 ————— 0s 554ms/step - accuracy: 0.7785 - loss: 0.6954  
Epoch 28: val\_accuracy improved from 0.80497 to 0.81883, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 95s 622ms/step - accuracy: 0.7785 - loss: 0.6955 - val\_accuracy: 0.8188 - val\_loss: 0.5974 - learning\_rate: 0.0010  
Epoch 29/100  
153/153 ————— 0s 553ms/step - accuracy: 0.7873 - loss: 0.6665  
Epoch 29: val\_accuracy did not improve from 0.81883  
153/153 ————— 94s 615ms/step - accuracy: 0.7873 - loss: 0.6666 - val\_accuracy: 0.8188 - val\_loss: 0.5843 - learning\_rate: 0.0010

Epoch 30/100

153/153 ————— 0s 554ms/step - accuracy: 0.7897 - loss: 0.6526

Epoch 30: val\_accuracy improved from 0.81883 to 0.82505, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras

153/153 ————— 95s 623ms/step - accuracy: 0.7897 - loss: 0.6525 - val\_accuracy: 0.8250 - val\_loss: 0.5622 - learning\_rate: 0.0010

Epoch 31/100

153/153 ————— 0s 554ms/step - accuracy: 0.7924 - loss: 0.6102

Epoch 31: val\_accuracy improved from 0.82505 to 0.84704, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras

153/153 ————— 95s 623ms/step - accuracy: 0.7925 - loss: 0.6102 - val\_accuracy: 0.8470 - val\_loss: 0.4839 - learning\_rate: 0.0010

Epoch 32/100

153/153 ————— 0s 554ms/step - accuracy: 0.8216 - loss: 0.5631

Epoch 32: val\_accuracy did not improve from 0.84704

153/153 ————— 94s 615ms/step - accuracy: 0.8216 - loss: 0.5633 - val\_accuracy: 0.8451 - val\_loss: 0.4960 - learning\_rate: 0.0010

Epoch 33/100

153/153 ————— 0s 557ms/step - accuracy: 0.8115 - loss: 0.5891

Epoch 33: val\_accuracy improved from 0.84704 to 0.85038, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras

153/153 ————— 96s 626ms/step - accuracy: 0.8115 - loss: 0.5891 - val\_accuracy: 0.8504 - val\_loss: 0.4976 - learning\_rate: 0.0010

Epoch 34/100

153/153 ————— 0s 555ms/step - accuracy: 0.8295 - loss: 0.5338

Epoch 34: val\_accuracy improved from 0.85038 to 0.85516, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras

153/153 ————— 95s 624ms/step - accuracy: 0.8295 - loss: 0.5339 - val\_accuracy: 0.8552 - val\_loss: 0.4705 - learning\_rate: 0.0010

Epoch 35/100

153/153 ————— 0s 555ms/step - accuracy: 0.8396 - loss: 0.5001

Epoch 35: val\_accuracy improved from 0.85516 to 0.87906, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras

153/153 ————— 95s 624ms/step - accuracy: 0.8396 - loss: 0.5002 - val\_accuracy: 0.8791 - val\_loss: 0.4209 - learning\_rate: 0.0010

Epoch 36/100

153/153 ————— 0s 555ms/step - accuracy: 0.8392 - loss: 0.5087

Epoch 36: val\_accuracy improved from 0.87906 to 0.88910, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras

153/153 ————— 95s 624ms/step - accuracy: 0.8392 - loss: 0.5087 - val\_accuracy: 0.8891 - val\_loss: 0.3902 - learning\_rate: 0.0010

Epoch 37/100

153/153 ————— 0s 559ms/step - accuracy: 0.8689 - loss: 0.4051

Epoch 37: val\_accuracy did not improve from 0.88910

153/153 ————— 95s 620ms/step - accuracy: 0.8687 - loss: 0.4055 - val\_accuracy: 0.8767 - val\_loss: 0.4262 - learning\_rate: 0.0010

Epoch 38/100

153/153 ————— 0s 555ms/step - accuracy: 0.8605 - loss: 0.4274

Epoch 38: val\_accuracy did not improve from 0.88910

153/153 ————— 94s 617ms/step - accuracy: 0.8605 - loss: 0.4276 - val\_accuracy: 0.8824 - val\_loss: 0.3937 - learning\_rate: 0.0010

Epoch 39/100

153/153 ————— 0s 556ms/step - accuracy: 0.8700 - loss: 0.4274

Epoch 39: val\_accuracy did not improve from 0.88910

153/153 ————— 94s 618ms/step - accuracy: 0.8700 - loss: 0.4276 - val\_accuracy: 0.8748 - val\_loss: 0.4118 - learning\_rate: 0.0010

Epoch 40/100  
153/153 ————— 0s 556ms/step - accuracy: 0.8626 - loss: 0.4486  
Epoch 40: val\_accuracy did not improve from 0.88910  
153/153 ————— 94s 617ms/step - accuracy: 0.8626 - loss: 0.4485 - val\_accuracy: 0.8853 - val\_loss: 0.3712 - learning\_rate: 0.0010  
Epoch 41/100  
153/153 ————— 0s 556ms/step - accuracy: 0.8583 - loss: 0.4358  
Epoch 41: val\_accuracy did not improve from 0.88910  
153/153 ————— 95s 618ms/step - accuracy: 0.8583 - loss: 0.4358 - val\_accuracy: 0.8881 - val\_loss: 0.3763 - learning\_rate: 0.0010  
Epoch 42/100  
153/153 ————— 0s 556ms/step - accuracy: 0.8797 - loss: 0.3651  
Epoch 42: val\_accuracy did not improve from 0.88910  
153/153 ————— 95s 619ms/step - accuracy: 0.8797 - loss: 0.3652 - val\_accuracy: 0.8748 - val\_loss: 0.4178 - learning\_rate: 0.0010  
Epoch 43/100  
153/153 ————— 0s 557ms/step - accuracy: 0.8894 - loss: 0.3435  
Epoch 43: val\_accuracy did not improve from 0.88910  
153/153 ————— 95s 619ms/step - accuracy: 0.8893 - loss: 0.3436 - val\_accuracy: 0.8891 - val\_loss: 0.3613 - learning\_rate: 0.0010  
Epoch 44/100  
153/153 ————— 0s 557ms/step - accuracy: 0.8754 - loss: 0.3861  
Epoch 44: val\_accuracy improved from 0.88910 to 0.89006, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 627ms/step - accuracy: 0.8754 - loss: 0.3862 - val\_accuracy: 0.8901 - val\_loss: 0.3779 - learning\_rate: 0.0010  
Epoch 45/100  
153/153 ————— 0s 557ms/step - accuracy: 0.8816 - loss: 0.3582  
Epoch 45: val\_accuracy improved from 0.89006 to 0.89293, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 628ms/step - accuracy: 0.8817 - loss: 0.3582 - val\_accuracy: 0.8929 - val\_loss: 0.3859 - learning\_rate: 0.0010  
Epoch 46/100  
153/153 ————— 0s 557ms/step - accuracy: 0.8991 - loss: 0.3057  
Epoch 46: val\_accuracy improved from 0.89293 to 0.90057, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 628ms/step - accuracy: 0.8991 - loss: 0.3059 - val\_accuracy: 0.9006 - val\_loss: 0.3758 - learning\_rate: 0.0010  
Epoch 47/100  
153/153 ————— 0s 557ms/step - accuracy: 0.9011 - loss: 0.2994  
Epoch 47: val\_accuracy did not improve from 0.90057  
153/153 ————— 95s 620ms/step - accuracy: 0.9010 - loss: 0.2995 - val\_accuracy: 0.8872 - val\_loss: 0.3851 - learning\_rate: 0.0010  
Epoch 48/100  
153/153 ————— 0s 557ms/step - accuracy: 0.8975 - loss: 0.3118  
Epoch 48: val\_accuracy improved from 0.90057 to 0.90296, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 628ms/step - accuracy: 0.8975 - loss: 0.3119 - val\_accuracy: 0.9030 - val\_loss: 0.3267 - learning\_rate: 0.0010  
Epoch 49/100  
153/153 ————— 0s 558ms/step - accuracy: 0.9076 - loss: 0.3046  
Epoch 49: val\_accuracy did not improve from 0.90296  
153/153 ————— 95s 621ms/step - accuracy: 0.9075 - loss: 0.3046 - val\_accuracy: 0.8963 - val\_loss: 0.3408 - learning\_rate: 0.0010  
Epoch 50/100  
153/153 ————— 0s 558ms/step - accuracy: 0.8965 - loss: 0.3341



Epoch 50: val\_accuracy did not improve from 0.90296  
153/153 ————— 95s 621ms/step - accuracy: 0.8966 - loss: 0.3340 - val\_accuracy: 0.9011 - val\_loss: 0.3504 - learning\_rate: 0.0010  
Epoch 51/100  
153/153 ————— 0s 558ms/step - accuracy: 0.8965 - loss: 0.3368  
Epoch 51: val\_accuracy did not improve from 0.90296  
153/153 ————— 95s 622ms/step - accuracy: 0.8965 - loss: 0.3367 - val\_accuracy: 0.8991 - val\_loss: 0.3620 - learning\_rate: 0.0010  
Epoch 52/100  
153/153 ————— 0s 558ms/step - accuracy: 0.9097 - loss: 0.2777  
Epoch 52: val\_accuracy improved from 0.90296 to 0.90344, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 629ms/step - accuracy: 0.9096 - loss: 0.2778 - val\_accuracy: 0.9034 - val\_loss: 0.3450 - learning\_rate: 0.0010  
Epoch 53/100  
153/153 ————— 0s 559ms/step - accuracy: 0.9050 - loss: 0.3126  
Epoch 53: val\_accuracy improved from 0.90344 to 0.90583, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 630ms/step - accuracy: 0.9050 - loss: 0.3126 - val\_accuracy: 0.9058 - val\_loss: 0.3324 - learning\_rate: 0.0010  
Epoch 54/100  
153/153 ————— 0s 560ms/step - accuracy: 0.9153 - loss: 0.2608  
Epoch 54: val\_accuracy did not improve from 0.90583  
153/153 ————— 95s 623ms/step - accuracy: 0.9153 - loss: 0.2609 - val\_accuracy: 0.9039 - val\_loss: 0.3469 - learning\_rate: 0.0010  
Epoch 55/100  
153/153 ————— 0s 559ms/step - accuracy: 0.9126 - loss: 0.2746  
Epoch 55: val\_accuracy did not improve from 0.90583  
153/153 ————— 95s 623ms/step - accuracy: 0.9126 - loss: 0.2746 - val\_accuracy: 0.9030 - val\_loss: 0.3496 - learning\_rate: 0.0010  
Epoch 56/100  
153/153 ————— 0s 559ms/step - accuracy: 0.9104 - loss: 0.2888  
Epoch 56: val\_accuracy did not improve from 0.90583  
153/153 ————— 95s 623ms/step - accuracy: 0.9104 - loss: 0.2887 - val\_accuracy: 0.9034 - val\_loss: 0.4116 - learning\_rate: 0.0010  
Epoch 57/100  
153/153 ————— 0s 559ms/step - accuracy: 0.9192 - loss: 0.2430  
Epoch 57: val\_accuracy improved from 0.90583 to 0.90631, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 97s 631ms/step - accuracy: 0.9191 - loss: 0.2431 - val\_accuracy: 0.9063 - val\_loss: 0.3559 - learning\_rate: 0.0010  
Epoch 58/100  
153/153 ————— 0s 560ms/step - accuracy: 0.9182 - loss: 0.2632  
Epoch 58: val\_accuracy improved from 0.90631 to 0.90822, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 631ms/step - accuracy: 0.9182 - loss: 0.2632 - val\_accuracy: 0.9082 - val\_loss: 0.3674 - learning\_rate: 0.0010  
Epoch 59/100  
153/153 ————— 0s 559ms/step - accuracy: 0.9219 - loss: 0.2478  
Epoch 59: val\_accuracy improved from 0.90822 to 0.91252, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras  
153/153 ————— 96s 631ms/step - accuracy: 0.9219 - loss: 0.2479 - val\_accuracy: 0.9125 - val\_loss: 0.3532 - learning\_rate: 0.0010  
Epoch 60/100  
153/153 ————— 0s 560ms/step - accuracy: 0.9249 - loss: 0.2306  
Epoch 60: val\_accuracy improved from 0.91252 to 0.92113, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras



```
min\Desktop\COSC2753_A2_MachineLearning/paddy_models/model_vgg_new.keras
153/153 ————— 97s 631ms/step - accuracy: 0.9249 - loss: 0.2307 - val_
accuracy: 0.9211 - val_loss: 0.3096 - learning_rate: 0.0010
Epoch 61/100
153/153 ————— 0s 559ms/step - accuracy: 0.9255 - loss: 0.2327
Epoch 61: val_accuracy did not improve from 0.92113
153/153 ————— 95s 624ms/step - accuracy: 0.9255 - loss: 0.2328 - val_
accuracy: 0.9101 - val_loss: 0.3489 - learning_rate: 0.0010
Epoch 62/100
153/153 ————— 0s 561ms/step - accuracy: 0.9118 - loss: 0.2709
Epoch 62: val_accuracy did not improve from 0.92113
153/153 ————— 96s 625ms/step - accuracy: 0.9118 - loss: 0.2709 - val_
accuracy: 0.9025 - val_loss: 0.4081 - learning_rate: 0.0010
Epoch 63/100
153/153 ————— 0s 560ms/step - accuracy: 0.9158 - loss: 0.2511
Epoch 63: val_accuracy did not improve from 0.92113
153/153 ————— 96s 625ms/step - accuracy: 0.9158 - loss: 0.2511 - val_
accuracy: 0.9144 - val_loss: 0.3552 - learning_rate: 0.0010
Epoch 64/100
153/153 ————— 0s 562ms/step - accuracy: 0.9277 - loss: 0.2288
Epoch 64: val_accuracy did not improve from 0.92113
153/153 ————— 96s 626ms/step - accuracy: 0.9277 - loss: 0.2289 - val_
accuracy: 0.9168 - val_loss: 0.3285 - learning_rate: 0.0010
Epoch 65/100
153/153 ————— 0s 562ms/step - accuracy: 0.9327 - loss: 0.2075
Epoch 65: val_accuracy did not improve from 0.92113
153/153 ————— 96s 626ms/step - accuracy: 0.9327 - loss: 0.2076 - val_
accuracy: 0.9125 - val_loss: 0.3608 - learning_rate: 0.0010
Epoch 66/100
153/153 ————— 0s 561ms/step - accuracy: 0.9211 - loss: 0.2482
Epoch 66: val_accuracy did not improve from 0.92113
153/153 ————— 96s 625ms/step - accuracy: 0.9211 - loss: 0.2482 - val_
accuracy: 0.9082 - val_loss: 0.3517 - learning_rate: 0.0010
Epoch 67/100
153/153 ————— 0s 562ms/step - accuracy: 0.9330 - loss: 0.2385
Epoch 67: val_accuracy did not improve from 0.92113
153/153 ————— 96s 627ms/step - accuracy: 0.9330 - loss: 0.2385 - val_
accuracy: 0.9015 - val_loss: 0.3830 - learning_rate: 0.0010
Epoch 68/100
153/153 ————— 0s 562ms/step - accuracy: 0.9333 - loss: 0.2283
Epoch 68: val_accuracy did not improve from 0.92113
153/153 ————— 96s 626ms/step - accuracy: 0.9333 - loss: 0.2283 - val_
accuracy: 0.9106 - val_loss: 0.3762 - learning_rate: 0.0010
Epoch 69/100
153/153 ————— 0s 562ms/step - accuracy: 0.9178 - loss: 0.2638
Epoch 69: val_accuracy did not improve from 0.92113
153/153 ————— 96s 626ms/step - accuracy: 0.9179 - loss: 0.2637 - val_
accuracy: 0.9192 - val_loss: 0.3676 - learning_rate: 0.0010
Epoch 70/100
153/153 ————— 0s 563ms/step - accuracy: 0.9420 - loss: 0.1798
Epoch 70: val_accuracy did not improve from 0.92113

Epoch 70: ReduceLROnPlateau reducing learning rate to 0.00010000000474974513.
153/153 ————— 96s 627ms/step - accuracy: 0.9419 - loss: 0.1800 - val_
accuracy: 0.9054 - val_loss: 0.3696 - learning_rate: 0.0010
Epoch 71/100
```

153/153 ————— 0s 563ms/step - accuracy: 0.9373 - loss: 0.1975  
Epoch 71: val\_accuracy did not improve from 0.92113

153/153 ————— 96s 628ms/step - accuracy: 0.9373 - loss: 0.1974 - val\_accuracy: 0.9183 - val\_loss: 0.3147 - learning\_rate: 1.0000e-04  
Epoch 72/100

153/153 ————— 0s 562ms/step - accuracy: 0.9566 - loss: 0.1397  
Epoch 72: val\_accuracy improved from 0.92113 to 0.92686, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras

153/153 ————— 97s 634ms/step - accuracy: 0.9566 - loss: 0.1397 - val\_accuracy: 0.9269 - val\_loss: 0.3151 - learning\_rate: 1.0000e-04  
Epoch 73/100

153/153 ————— 0s 562ms/step - accuracy: 0.9600 - loss: 0.1290  
Epoch 73: val\_accuracy did not improve from 0.92686

153/153 ————— 96s 627ms/step - accuracy: 0.9600 - loss: 0.1290 - val\_accuracy: 0.9269 - val\_loss: 0.3121 - learning\_rate: 1.0000e-04  
Epoch 74/100

153/153 ————— 0s 563ms/step - accuracy: 0.9635 - loss: 0.1170  
Epoch 74: val\_accuracy improved from 0.92686 to 0.93308, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras

153/153 ————— 97s 636ms/step - accuracy: 0.9635 - loss: 0.1169 - val\_accuracy: 0.9331 - val\_loss: 0.3123 - learning\_rate: 1.0000e-04  
Epoch 75/100

153/153 ————— 0s 564ms/step - accuracy: 0.9615 - loss: 0.1066  
Epoch 75: val\_accuracy did not improve from 0.93308

153/153 ————— 96s 629ms/step - accuracy: 0.9615 - loss: 0.1066 - val\_accuracy: 0.9321 - val\_loss: 0.3177 - learning\_rate: 1.0000e-04  
Epoch 76/100

153/153 ————— 0s 561ms/step - accuracy: 0.9680 - loss: 0.1010  
Epoch 76: val\_accuracy did not improve from 0.93308

153/153 ————— 96s 626ms/step - accuracy: 0.9679 - loss: 0.1010 - val\_accuracy: 0.9307 - val\_loss: 0.3128 - learning\_rate: 1.0000e-04  
Epoch 77/100

153/153 ————— 0s 562ms/step - accuracy: 0.9732 - loss: 0.0791  
Epoch 77: val\_accuracy improved from 0.93308 to 0.93356, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras

153/153 ————— 97s 634ms/step - accuracy: 0.9732 - loss: 0.0791 - val\_accuracy: 0.9336 - val\_loss: 0.3135 - learning\_rate: 1.0000e-04  
Epoch 78/100

153/153 ————— 0s 562ms/step - accuracy: 0.9677 - loss: 0.0974  
Epoch 78: val\_accuracy did not improve from 0.93356

153/153 ————— 96s 627ms/step - accuracy: 0.9677 - loss: 0.0974 - val\_accuracy: 0.9316 - val\_loss: 0.3194 - learning\_rate: 1.0000e-04  
Epoch 79/100

153/153 ————— 0s 561ms/step - accuracy: 0.9690 - loss: 0.0968  
Epoch 79: val\_accuracy improved from 0.93356 to 0.93451, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras

153/153 ————— 97s 634ms/step - accuracy: 0.9690 - loss: 0.0967 - val\_accuracy: 0.9345 - val\_loss: 0.3167 - learning\_rate: 1.0000e-04  
Epoch 80/100

153/153 ————— 0s 562ms/step - accuracy: 0.9733 - loss: 0.0891  
Epoch 80: val\_accuracy improved from 0.93451 to 0.93690, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning/paddy\_models/model\_vgg\_new.keras

153/153 ————— 97s 634ms/step - accuracy: 0.9733 - loss: 0.0892 - val\_accuracy: 0.9369 - val\_loss: 0.3250 - learning\_rate: 1.0000e-04  
Epoch 81/100

153/153 ————— 0s 562ms/step - accuracy: 0.9746 - loss: 0.0756


Epoch 81: val\_accuracy did not improve from 0.93690  
153/153 ————— 96s 626ms/step - accuracy: 0.9746 - loss: 0.0756 - val\_accuracy: 0.9364 - val\_loss: 0.3149 - learning\_rate: 1.0000e-04  
Epoch 82/100  
153/153 ————— 0s 562ms/step - accuracy: 0.9735 - loss: 0.0875  
Epoch 82: val\_accuracy did not improve from 0.93690  
153/153 ————— 96s 626ms/step - accuracy: 0.9735 - loss: 0.0875 - val\_accuracy: 0.9350 - val\_loss: 0.3172 - learning\_rate: 1.0000e-04  
Epoch 83/100  
153/153 ————— 0s 563ms/step - accuracy: 0.9734 - loss: 0.0735  
Epoch 83: val\_accuracy did not improve from 0.93690  
153/153 ————— 96s 627ms/step - accuracy: 0.9734 - loss: 0.0736 - val\_accuracy: 0.9336 - val\_loss: 0.3157 - learning\_rate: 1.0000e-04  
Epoch 84/100  
153/153 ————— 0s 561ms/step - accuracy: 0.9774 - loss: 0.0642  
Epoch 84: val\_accuracy did not improve from 0.93690  
153/153 ————— 96s 626ms/step - accuracy: 0.9774 - loss: 0.0642 - val\_accuracy: 0.9345 - val\_loss: 0.3348 - learning\_rate: 1.0000e-04  
Epoch 85/100  
153/153 ————— 0s 562ms/step - accuracy: 0.9793 - loss: 0.0654  
Epoch 85: val\_accuracy did not improve from 0.93690  
153/153 ————— 96s 627ms/step - accuracy: 0.9793 - loss: 0.0655 - val\_accuracy: 0.9359 - val\_loss: 0.3188 - learning\_rate: 1.0000e-04  
Epoch 86/100  
153/153 ————— 0s 561ms/step - accuracy: 0.9718 - loss: 0.0838  
Epoch 86: val\_accuracy did not improve from 0.93690  
153/153 ————— 96s 626ms/step - accuracy: 0.9718 - loss: 0.0838 - val\_accuracy: 0.9359 - val\_loss: 0.3296 - learning\_rate: 1.0000e-04  
Epoch 87/100  
153/153 ————— 0s 564ms/step - accuracy: 0.9725 - loss: 0.0756  
Epoch 87: val\_accuracy did not improve from 0.93690  
153/153 ————— 96s 629ms/step - accuracy: 0.9725 - loss: 0.0756 - val\_accuracy: 0.9355 - val\_loss: 0.3255 - learning\_rate: 1.0000e-04  
Epoch 88/100  
153/153 ————— 0s 562ms/step - accuracy: 0.9731 - loss: 0.0816  
Epoch 88: val\_accuracy did not improve from 0.93690  
153/153 ————— 96s 626ms/step - accuracy: 0.9731 - loss: 0.0816 - val\_accuracy: 0.9364 - val\_loss: 0.3304 - learning\_rate: 1.0000e-04  
Epoch 89/100  
153/153 ————— 0s 563ms/step - accuracy: 0.9751 - loss: 0.0648  
Epoch 89: val\_accuracy did not improve from 0.93690  
153/153 ————— 96s 627ms/step - accuracy: 0.9751 - loss: 0.0648 - val\_accuracy: 0.9345 - val\_loss: 0.3354 - learning\_rate: 1.0000e-04  
Epoch 90/100  
153/153 ————— 0s 561ms/step - accuracy: 0.9765 - loss: 0.0734  
Epoch 90: val\_accuracy improved from 0.93690 to 0.93881, saving model to c:\Users\Admin\Desktop\COSC2753\_A2\_MachineLearning\paddy\_models\model\_vgg\_new.keras  
153/153 ————— 97s 633ms/step - accuracy: 0.9765 - loss: 0.0734 - val\_accuracy: 0.9388 - val\_loss: 0.3308 - learning\_rate: 1.0000e-04  
Epoch 91/100  
153/153 ————— 0s 561ms/step - accuracy: 0.9767 - loss: 0.0631  
Epoch 91: val\_accuracy did not improve from 0.93881  
153/153 ————— 96s 626ms/step - accuracy: 0.9767 - loss: 0.0631 - val\_accuracy: 0.9379 - val\_loss: 0.3335 - learning\_rate: 1.0000e-04  
Epoch 92/100  
153/153 ————— 0s 563ms/step - accuracy: 0.9775 - loss: 0.0726

Epoch 92: val\_accuracy did not improve from 0.93881  
**153/153** ————— **96s** 627ms/step - accuracy: 0.9775 - loss: 0.0725 - val\_accuracy: 0.9379 - val\_loss: 0.3388 - learning\_rate: 1.0000e-04  
Epoch 93/100  
**153/153** ————— **0s** 562ms/step - accuracy: 0.9759 - loss: 0.0695  
Epoch 93: val\_accuracy did not improve from 0.93881  
**153/153** ————— **96s** 627ms/step - accuracy: 0.9759 - loss: 0.0695 - val\_accuracy: 0.9379 - val\_loss: 0.3360 - learning\_rate: 1.0000e-04  
Epoch 94/100  
**153/153** ————— **0s** 563ms/step - accuracy: 0.9768 - loss: 0.0732  
Epoch 94: val\_accuracy did not improve from 0.93881  
**153/153** ————— **96s** 627ms/step - accuracy: 0.9768 - loss: 0.0732 - val\_accuracy: 0.9388 - val\_loss: 0.3174 - learning\_rate: 1.0000e-04  
Epoch 95/100  
**153/153** ————— **0s** 561ms/step - accuracy: 0.9808 - loss: 0.0669  
Epoch 95: val\_accuracy did not improve from 0.93881  
**153/153** ————— **96s** 626ms/step - accuracy: 0.9808 - loss: 0.0669 - val\_accuracy: 0.9359 - val\_loss: 0.3311 - learning\_rate: 1.0000e-04  
Epoch 96/100  
**153/153** ————— **0s** 583ms/step - accuracy: 0.9778 - loss: 0.0622  
Epoch 96: val\_accuracy did not improve from 0.93881  
**153/153** ————— **99s** 648ms/step - accuracy: 0.9778 - loss: 0.0622 - val\_accuracy: 0.9369 - val\_loss: 0.3416 - learning\_rate: 1.0000e-04  
Epoch 97/100  
**153/153** ————— **0s** 575ms/step - accuracy: 0.9793 - loss: 0.0703  
Epoch 97: val\_accuracy did not improve from 0.93881  
**153/153** ————— **98s** 641ms/step - accuracy: 0.9793 - loss: 0.0703 - val\_accuracy: 0.9355 - val\_loss: 0.3371 - learning\_rate: 1.0000e-04  
Epoch 98/100  
**153/153** ————— **0s** 576ms/step - accuracy: 0.9823 - loss: 0.0465  
Epoch 98: val\_accuracy did not improve from 0.93881  
**153/153** ————— **98s** 642ms/step - accuracy: 0.9823 - loss: 0.0465 - val\_accuracy: 0.9340 - val\_loss: 0.3429 - learning\_rate: 1.0000e-04  
Epoch 99/100  
**153/153** ————— **0s** 577ms/step - accuracy: 0.9778 - loss: 0.0649  
Epoch 99: val\_accuracy did not improve from 0.93881  
**153/153** ————— **98s** 643ms/step - accuracy: 0.9778 - loss: 0.0649 - val\_accuracy: 0.9355 - val\_loss: 0.3320 - learning\_rate: 1.0000e-04  
Epoch 100/100  
**153/153** ————— **0s** 579ms/step - accuracy: 0.9832 - loss: 0.0519  
Epoch 100: val\_accuracy did not improve from 0.93881  
  
Epoch 100: ReduceLROnPlateau reducing learning rate to 1.0000000474974514e-05.  
**153/153** ————— **99s** 645ms/step - accuracy: 0.9832 - loss: 0.0520 - val\_accuracy: 0.9374 - val\_loss: 0.3330 - learning\_rate: 1.0000e-04  
**108/108** ————— **17s** 161ms/step - accuracy: 0.9255 - loss: 0.3696  
Test accuracy: 92.55%

```
In [51]: vit_classifier.save_weights(HOME_PATH + '/paddy_models/vit.weights.h5')
         vit_classifier.save(HOME_PATH + '/paddy_models/vit_model.keras')
```


```
In [43]: history = history
         loss = history.history['loss']
         val_loss = history.history['val_loss']
         epochs = range(1, len(loss) + 1)
         plt.plot(epochs, loss, 'y', label='Training loss')
```

```
plt.plot(epochs, val_loss, 'r', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
In [44]: acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.plot(epochs, acc, 'y', label='Training accuracy')
plt.plot(epochs, val_acc, 'r', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('accuracy')
plt.legend()
plt.show()
```



## Generate Prediction

```
In [52]: validation_ds = keras.utils.image_dataset_from_directory(
    directory=HOME_PATH + 'train_images',
    batch_size=16,
    image_size=(256, 256),
    validation_split=0.2,
    subset="validation",
    seed=123
)



test_ds = keras.utils.image_dataset_from_directory(
    directory = HOME_PATH + 'test_images',
    batch_size = 16,
    image_size = (256, 256),
    label_mode = None,
    shuffle=False
)

# Predict the labels of the test set
predictions = vit_classifier.predict(test_ds)

predicted_labels = [labels[prediction.argmax()] for prediction in predictions]

loss, accuracy = vit_classifier.evaluate(validation_ds)
print(f'Validation Loss: {loss:.4f}, Validation Accuracy: {accuracy:.4f}')
```



Found 10407 files belonging to 10 classes.  
Using 2081 files for validation.  
Found 3469 files.  
217/217  19s 86ms/step  
131/131  11s 84ms/step - accuracy: 0.9603 - loss: 0.1769  
Validation Loss: 0.1875, Validation Accuracy: 0.9611

```
In [53]: # Create a submission file
submission_df = pd.DataFrame({'image_id': test_ds.file_paths, 'label': predicted_la

submission_df['image_id'] = submission_df['image_id'].apply(lambda x: x.split('/')[

submission_df.to_csv('sample_submission.csv', index=False)
```

```
In [54]: # Display unique predicted labels
print(set(predicted_labels))

# Display submission head
print(submission_df.head())
```

```
{np.int64(0)}
      image_id  label
0  test_images\200001.jpg      0
1  test_images\200002.jpg      0
2  test_images\200003.jpg      0
3  test_images\200004.jpg      0
4  test_images\200005.jpg      0
```